

The package {bigstatsr}: memory- and computation-efficient tools for big matrices stored on disk

Florian Privé (@privefl)

Rencontres R 2018

Slides: <https://privefl.github.io/RR18/bigstatsr.html>

Motivation


Analyze very large genotype matrices

- previously: 15K x 280K, **celiac disease** (~30GB)
- currently: 500K x 500K, **UK Biobank** (~2TB)



But I still want to use **R**..

The solution I found



Format FBM is very similar to format `filebacked.big.matrix` from package `{bigmemory}` (details in [this vignette](#)).

Simple accessors

Similar accessor as R matrices

```
X <- FBM(2, 5, init = 1:10, backingfile = "test")
```

```
X$backingfile
```

```
## [1] "/home/privef/Bureau/RR18/test.bk"
```

```
X[, 1] ## ok
```

```
## [1] 1 2
```

```
X[1, ] ## bad
```

```
## [1] 1 3 5 7 9
```

```
X[] ## super bad
```

```
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]    1    3    5    7    9  
## [2,]    2    4    6    8   10
```

Similar accessor as R matrices


```
colSums(X[])  ## super bad
```

```
## [1] 3 7 11 15 19
```



Split-(par)Apply-Combine Strategy

Apply standard R functions to big matrices (in parallel)



Implemented in `big_apply()`.

Similar accessor as Rcpp matrices

```
// [[Rcpp::depends(BH, bigstatsr)]]
#include <bigstatsr/BMAcc.h>

// [[Rcpp::export]]
NumericVector big_colsums(Environment BM) {

  XPtr<FBM> xpBM = BM["address"];
  BMAcc<double> macc(xpBM);

  size_t n = macc.nrow();
  size_t m = macc.ncol();

  NumericVector res(m);


  for (size_t j = 0; j < m; j++)
    for (size_t i = 0; i < n; i++)
      res[j] += macc(i, j);

  return res;
}
```

Some examples
from my work

Partial Singular Value Decomposition

15K \times 100K -- 10 first PCs -- 6 cores -- **1 min** (vs 2h in base R)




Implemented in `big_randomSVD()`, powered by R packages `{RSpectra}` and `{Rcpp}`.

Sparse linear models


Predicting complex diseases via penalized logistic regression

15K \times 280K -- 6 cores -- 2 min



Multiple association testing

Which DNA mutations are associated with one disease?



Conclusion

I'm able to run algorithms
on 100GB of data
in  on my computer

Advantages of using FBM objects

- you can apply algorithms on **data larger than your RAM**,
- you can easily **parallelize** your algorithms because the data on disk is shared,
- you write **more efficient algorithms** (you do less copies and think more about what you're doing),
- you can use **different types of data**, for example, in my field, I'm storing my data with only 1 byte per element (rather than 8 bytes for a standard R matrix). See [the documentation of the FBM class](#) for details.

R Packages

Efficient analysis of large-scale genome-wide data with two R packages: bigstatsr and bigsnpr

Florian Privé , Hugues Aschard, Andrey Ziyatdinov, Michael G B Blum 

Bioinformatics, bty185, <https://doi.org/10.1093/bioinformatics/bty185>

- {bigstatsr}: to be used by any field of research
- {bigsnpr}: algorithms specific to my field of research

Contributors are welcomed!



Make sure to grab an hex sticker



Thanks!

Presentation: <https://privefl.github.io/RR18/bigstatsr.html>

Package's website: <https://privefl.github.io/bigstatsr/>

DOI: [10.1093/bioinformatics/bty185](https://doi.org/10.1093/bioinformatics/bty185)



privefl



privefl



F. Privé

Slides created via the R package **xaringan**.